

Reproducibility Study and Inference-Time Optimization of Denoising Diffusion Probabilistic Models on CIFAR-10

Zain Shahid

Dept. of Data Science
National University of Computer
and Emerging Sciences, Islamabad
Islamabad, Pakistan
i232582@isb.nu.edu.pk

Muhammad Talha Arshad

Dept. of Data Science
National University of Computer
and Emerging Sciences, Islamabad
Islamabad, Pakistan
i232548@isb.nu.edu.pk

Sanaullah Farooqi

Dept. of Data Science
National University of Computer
and Emerging Sciences, Islamabad
Islamabad, Pakistan
i232594@isb.nu.edu.pk

Abstract—Denoising Diffusion Probabilistic Models (DDPM), introduced by Ho et al., represent a foundational class of generative models that achieve high-quality image synthesis through iterative denoising. This paper presents a complete reproducibility study and inference-time optimization analysis of DDPM on the CIFAR-10 benchmark, conducted as part of a three-part deep learning course project at FAST-NUCES Islamabad. The work involved overcoming significant real-world challenges: GPU architecture incompatibility requiring two days of debugging before resolving with a PyTorch nightly build, permanently deleted official model checkpoints preventing reproduction of baseline results, power interruptions during multi-hour generation runs requiring process recovery and corruption checking across 15,000 images, accidental simultaneous duplicate process launches causing write race conditions, and unavailability of FID reference statistics for the cross-domain evaluation dataset. Despite these obstacles, we successfully reproduced the DDPM pipeline using the official pretrained checkpoint and the HuggingFace diffusers library, achieving a Frechet Inception Distance (FID) of 20.91 against the paper-reported 3.17, generating 15,000 images across three sampling configurations. Building on this, we propose and conduct three original experiments: (1) a DDIM step count ablation across nine configurations from 1 to 1000 steps identifying an optimal at 200 steps with FID 42.13, where 50 steps achieves 94 percent of optimal quality at 20 percent of computation cost; (2) an eta parameter study at fixed 10 steps identifying a phase-transition degradation above eta 0.5, where deterministic sampling achieves FID 65.72 versus 89.55 for full stochastic, a 36 percent gap; and (3) a cross-domain CelebA-HQ 256x256 evaluation generating 244 face images over 7 hours reporting intra-FID 59.34. All code, logs, and reports are publicly available at <https://github.com/Minato-sudo/deep-learning-assignment2>.

Index Terms—Diffusion models, DDPM, DDIM, FID, image generation, CIFAR-10, inference optimization, reproducibility, CelebA-HQ, deep learning

I. INTRODUCTION

Generative models have become central to modern computer vision, enabling the synthesis of photorealistic images from learned data distributions. Among them, Denoising Diffusion Probabilistic Models (DDPM), introduced by Ho, Jain, and Abbeel [1], have established state-of-the-art results on standard

image generation benchmarks. On unconditional CIFAR-10 image generation at 32x32 resolution, DDPM achieves a Frechet Inception Distance (FID) of 3.17 and an Inception Score of 9.46, surpassing previous generative approaches including Generative Adversarial Networks in terms of both image quality and output diversity.

The core mechanism of DDPM involves a fixed forward process that gradually corrupts real images with Gaussian noise over T equal to 1000 timesteps according to a fixed variance schedule, and a learned reverse process parameterized by a U-Net neural network that recovers clean images by iteratively denoising from pure random noise. Despite impressive generation quality, the requirement for 1000 sequential neural network evaluations at inference time makes DDPM computationally expensive for practical deployment. This has motivated significant research into faster sampling strategies, most notably DDIM [2], which can generate images in as few as 10 to 50 steps.

This paper presents the complete record of our three-part course project. Assignment 1 focused on understanding the DDPM paper and related work. Assignment 2 focused on reproducing the DDPM results on CIFAR-10. Assignment 3 extended the work with three original experiments on inference-time hyperparameter optimization and cross-domain generalization. The project was conducted over several weeks on a student laptop equipped with an NVIDIA RTX 5050 GPU and involved overcoming a series of real technical and infrastructure challenges that are documented in full as part of this report. Honest documentation of these challenges, including hardware incompatibilities, deleted checkpoints, data corruption from power failures, and unavailable evaluation statistics, is itself a contribution to reproducibility science.

Our key contributions are as follows. We provide a fully documented reproduction of DDPM on CIFAR-10 including all setup steps, hardware challenges, failed attempts, and corrective actions taken, together with honest quantitative analysis of the FID gap. We empirically identify the optimal DDIM step count as 200 steps and a practical efficient configuration at 50

steps achieving 94 percent of peak quality. We characterize a previously undescribed phase transition in DDIM image quality at eta equal to 0.5 under a 10-step budget, providing concrete deployment recommendations. We provide a cross-domain evaluation of pre-trained DDPM weights on human face synthesis at 256x256 resolution under real computational constraints requiring 7 hours of continuous GPU operation.

II. BACKGROUND AND RELATED WORK

A. Assignment 1 Paper Understanding: DDPM

In Assignment 1 we performed a thorough reading and analysis of the primary paper, Denoising Diffusion Probabilistic Models by Ho, Jain, and Abbeel [1]. We identified the core mathematical framework, the U-Net architecture used as the noise predictor, the training objective, and the evaluation protocol. We also studied two baseline papers: Consistency Models [3] and EDM [4]. The key insight from Assignment 1 was that DDPM’s quality comes from the iterative denoising process over many steps, and that the model is stable to train compared to GANs because it uses a simple mean-squared error objective rather than an adversarial loss.

B. Denoising Diffusion Probabilistic Models

Ho, Jain, and Abbeel [1] formalized DDPM as a latent variable model inspired by non-equilibrium thermodynamics [9]. The forward process adds Gaussian noise to real data according to a fixed variance schedule over T timesteps:

$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I}) \quad (1)$$

The reverse process is a learned Markov chain parameterized by a U-Net trained to predict the noise component using the simplified mean-squared error objective:

$$L_{\text{simple}} = \mathbb{E}_{t, \mathbf{x}_0, \epsilon} [\|\epsilon - \epsilon_\theta(\mathbf{x}_t, t)\|^2] \quad (2)$$

On unconditional CIFAR-10, the model achieves FID 3.17 and Inception Score 9.46 using 50,000 generated samples evaluated against all 50,000 training images. The official TensorFlow implementation is available at <https://github.com/hojonathanho/diffusion> and the pretrained weights are accessible via HuggingFace as google/ddpm-cifar10-32.

C. Denoising Diffusion Implicit Models (DDIM)

Song, Meng, and Ermon [2] introduced DDIM as a non-Markovian generalization of the DDPM reverse process. DDIM defines a deterministic sampling trajectory parameterized by a subset of timesteps, enabling high-quality generation in 10 to 50 steps without retraining. The key parameter eta controls stochasticity: eta equal to 0 gives fully deterministic ODE-based sampling, while eta equal to 1 recovers DDPM-equivalent stochastic sampling.

D. EDM: Design Space of Diffusion Models

Karras, Aittala, Aila, and Laine [4] analyzed diffusion model design choices including preconditioning, noise schedules, and ODE solver order. By introducing a higher-order Heun sampler and improved preconditioning of the score network, EDM achieves FID 1.97 (unconditional) and 1.79 (class-conditional) on CIFAR-10 using only 35 network evaluations. Their work established that first-order Euler solvers accumulate discretization errors, which is directly relevant to our finding that DDIM degrades at 1000 steps.

E. Consistency Models

Song, Dhariwal, Chen, and Sutskever [3] proposed Consistency Models for single-step generation by enforcing self-consistency along the probability flow ODE. They achieve FID 3.55 in one step and 2.93 in two steps on CIFAR-10. The official checkpoint was permanently deleted from OpenAI blob storage and HuggingFace before our evaluation, preventing direct reproduction.

F. Improved DDPM and Latent Diffusion

Nichol and Dhariwal [5] improved DDPM with learned variance schedules and cosine noise scheduling, achieving FID 2.90 on CIFAR-10. Rombach et al. [6] introduced Latent Diffusion Models (LDM), which apply diffusion in compressed latent space enabling high-resolution synthesis. SDXL [8] further extended LDM to multi-GPU high-resolution generation. We had initially proposed SDXL as a baseline but found it infeasible on a single laptop GPU and replaced it with DDPM as our primary paper.

III. PROPOSED MODEL AND EXPERIMENTAL DESIGN

A. Reproduction Architecture and Approach

Our reproduction uses the official pretrained checkpoint google/ddpm-cifar10-32 from HuggingFace, loaded via DDPM Pipeline in the diffusers library [10]. The U-Net architecture includes the following components:

- Encoder-decoder structure with symmetric skip connections at each resolution
- Residual blocks with group normalization at each spatial scale
- Self-attention layers at 16x16 spatial resolution for global context modeling
- Sinusoidal timestep positional embeddings added to each residual block
- Linear noise schedule from beta 0.0001 to 0.02 over T equal to 1000 timesteps
- Input and output: 32x32x3 for CIFAR-10

For DDIM experiments we switch to DDIM Pipeline with the same U-Net weights, varying only num_inference_steps and eta. This enables direct comparison of sampling strategies without any retraining.

B. Proposed Experiment 1: DDIM Step Count Ablation

Hypothesis: There exists an optimal number of inference steps beyond which quality saturates or degrades due to discretization errors in the first-order Euler ODE solver. We hypothesize the optimal lies between 50 and 200 steps based on prior literature.

Motivation: Assignment 2 tested only three step counts: 1, 10, and 1000. A fine-grained sweep across nine configurations fully characterizes the quality-speed tradeoff curve, enabling practical deployment recommendations. This experiment directly addresses the efficiency improvement goal of Assignment 3.

C. Proposed Experiment 2: DDIM Eta Parameter Study

Hypothesis: Deterministic sampling (eta equal to 0) outperforms stochastic sampling at low step budgets because added noise at each step compounds errors when the step count is small. Quality degrades monotonically as eta increases from 0 to 1.

Motivation: While DDIM theorizes a quality-diversity tradeoff controlled by eta, the practical effect at a fixed 10-step budget on our specific hardware and checkpoint had not been characterized. This experiment isolates eta as the sole variable while holding steps fixed at 10.

D. Proposed Experiment 3: Cross-Domain Generalization

Hypothesis: Pre-trained DDPM weights encode a general denoising prior. A separate DDPM checkpoint trained on human face images will produce realistic and diverse outputs, confirming that the framework generalizes across image domains.

Motivation: Assignment 3 requires at least one additional dataset. Cross-domain evaluation tests whether the DDPM framework is domain-agnostic, which has practical implications for applying pre-trained models to new image types. The original DDPM paper also reported results on LSUN 256x256, confirming the architecture scales to this resolution.

IV. EXPERIMENTAL SETUP AND CHALLENGES ENCOUNTERED

A. Hardware and Software Configuration

TABLE I
HARDWARE AND SOFTWARE CONFIGURATION

Component	Specification
GPU	NVIDIA GeForce RTX 5050 Laptop 8GB VRAM, sm_120 Blackwell arch.
CPU	Intel Core i7, 16 threads
RAM	16GB
OS	Kubuntu Linux (Ubuntu 24.04)
PyTorch	2.12.0.dev20260328+cu128 (nightly)
Python	3.12.3
diffusers	0.33.1
clean-fid	0.1.35
torchvision	0.20.0

B. Challenge 1: GPU Architecture Incompatibility (2 Days Lost)

The NVIDIA RTX 5050 uses the Blackwell architecture (compute capability sm_120), released in late 2024. Stable PyTorch did not support sm_120 and raised a fatal CUDA kernel error on any GPU tensor operation. We spent approximately two full days trying: reinstalling CUDA drivers, testing older PyTorch 2.x stable releases, attempting CPU-only fallback, and researching community forums. The correct solution was a PyTorch nightly build compiled against CUDA 12.8:

```
pip install --pre torch torchvision \
  --index-url https://download.pytorch.org/ \
  whl/nightly/cu128
```

After this, torch.cuda.is_available() returned True and all generation and FID computation ran on GPU.

C. Challenge 2: Baseline Paper Reassignment (3 Days Lost)

Our original Assignment 1 proposal included SDXL and a survey paper. Upon starting Assignment 2 we discovered SDXL requires multi-GPU data center hardware and the survey paper had no implementable model. We communicated this to our instructor Dr. Zohair Ahmed who acknowledged our submission. We replaced baselines with EDM and Consistency Models, then had to re-read both papers before restarting implementation. This added approximately three days to the project timeline.

D. Challenge 3: Deleted Official Checkpoints (3 Days Spent)

The Consistency Models checkpoint on OpenAI blob storage and HuggingFace returned HTTP 404 on all download attempts across three days. We tried the HuggingFace snapshot download API, the CLI tool, direct URL access, and archived versions. All failed. This made direct FID reproduction of Consistency Models impossible. We documented this in Assignment 2 and used DDIM step variations as a conceptual substitute. This situation reflects a real problem in deep learning reproducibility.

E. Challenge 4: Power Interruption and Data Recovery (3 Hours)

During the generation of 5000 DDPM images, a power fluctuation shut down the laptop, killing both background processes with 4500 images already written. We recovered by resuming both processes using kill -CONT on their PIDs. After completion we ran a systematic PIL Image verify check across all 15,000 images and found 5 corrupted files, which were deleted. Final counts confirmed as 5000 per folder.

F. Challenge 5: Duplicate Process Race Condition

We accidentally launched two instances of generate_more.py simultaneously, both writing to the same output folder. We detected this via ps aux showing two PIDs for the same script. We allowed both to complete, accepting small

corruption risk that PIL verify confirmed was only 5 files. For all Assignment 3 experiments we strictly checked ps aux before launching any script.

G. Challenge 6: CelebA-HQ Generation Time (7 Hours Continuous)

At 256x256 resolution, 1000 steps, and batch size 4, generation required approximately 90 seconds per image on our GPU. We ran continuously for over 7 hours to generate 244 images before stopping due to project deadline constraints. The original target was 500 images.

H. Challenge 7: FID Reference Statistics Unavailable (HTTP 404)

clean-fid reference statistics for CelebA-HQ 256x256 returned HTTP 404 during evaluation. We instead computed intra-FID between two equal halves of the 244 generated images (122 each) as a diversity metric, and cited published FID ranges from literature as qualitative context.

I. Datasets and Evaluation Protocol

CIFAR-10: 10,000 real CIFAR-10 images from torchvision saved as PNG files serve as the FID reference set. Generated images: 5,000 per configuration for Assignment 2, 1,000 per configuration for Assignment 3. Total images generated across the full project: over 24,000.

CelebA-HQ 256x256: 244 images generated using google/ddpm-celebahq-256 over 7 hours.

FID computed using clean-fid [11] with InceptionV3 features. Intra-FID computed by splitting 244 images into two groups of 122. Sample grids generated for qualitative evaluation at each configuration.

J. Hyperparameters and Configurations

TABLE II
ASSIGNMENT 2 GENERATION CONFIGURATIONS

Model	Steps	Eta	Samples
DDPM	1000	1.0	5000
DDIM	10	0.0	5000
DDIM	1	0.0	5000

TABLE III
ASSIGNMENT 3 EXPERIMENT CONFIGURATIONS

Experiment	Steps	Eta	Samples
Exp1 Step Ablation	1–1000 (9 configs)	0.0 fixed	1000 each
Exp2 Eta Study	10 fixed	0.0–1.0 (5)	1000 each
Exp3 Cross-Domain	1000	1.0	244

V. RESULTS

A. Assignment 2: Reproduced Results on CIFAR-10

Table IV presents our reproduction results using 5,000 generated samples per configuration.

The FID gap between our reproduced DDPM (20.91) and the paper (3.17) is attributed to: (1) sample size difference

TABLE IV
ASSIGNMENT 2 REPRODUCED RESULTS: CIFAR-10 (5000 SAMPLES)

Model	Steps	Paper FID	Our FID
DDPM	1000	3.17	20.91
DDIM (eta=0)	10	N/A	38.69
DDIM (eta=0)	1	N/A	462.11

— 5,000 versus 50,000 generated images where FID variance is inversely proportional to sample count; (2) reference set difference — 10,000 versus 50,000 real images; and (3) evaluation protocol differences in image preprocessing. The trend is fully consistent: DDPM at 1000 steps substantially outperforms DDIM at fewer steps, and DDIM at 1 step produces unusable noise.




Fig. 1. Real CIFAR-10 reference images (5x5 grid, 32x32 pixels). These 10,000 real images formed the FID reference distribution for all CIFAR-10 experiments.




Fig. 2. Assignment 2 FID comparison chart (log scale). Blue: our reproduced results with 5k samples. Green: paper-reported FID of 3.17. The gap is explained by sample size and evaluation protocol differences. The trend matches the paper perfectly.

B. Assignment 3 Experiment 1: DDIM Step Count Ablation

Finding 1 – Optimal at 200 steps: The best FID (42.13) occurs at 200 steps. Increasing steps beyond 200 actively degrades quality for deterministic DDIM with eta equal to 0.

Finding 2 – Degradation at 1000 steps (44 percent): FID rises from 42.13 at 200 steps to 60.55 at 1000 steps. This is




Fig. 3. DDPM generated images at 1000 steps (FID 20.91). Recognizable animals and vehicles confirm the model successfully learned the CIFAR-10 distribution.




Fig. 4. DDIM generated images at 10 steps, $\eta=0$ (FID 38.69). Visibly reduced sharpness and detail compared to 1000-step DDPM, consistent with the higher FID.



Fig. 5. DDIM generated images at 1 step, $\eta=0$ (FID 462.11). Completely unstructured random noise. A single denoising step is entirely insufficient to recover any image structure from pure noise.

TABLE V
EXPERIMENT 1: FID vs. INFERENCE STEPS (CIFAR-10, $\eta=0$, 1000 SAMPLES)

Steps	FID Score	vs. Best
1	468.11	+1011%
5	95.64	+127%
10	64.07	+52%
20	51.94	+23%
50	45.11	+7%
100	42.74	+1%
200	42.13	Best
500	44.02	+4%
1000	60.55	+44%




Fig. 6. FID vs. DDIM inference steps (log x-axis). Optimal at 200 steps (FID 42.13) marked by red dotted line. The green dashed line shows the paper-reported FID of 3.17. FID degrades at both extremes due to insufficient denoising (left) and ODE discretization error accumulation (right).

caused by accumulated first-order Euler solver discretization errors over many steps, consistent with Karras et al. [4].

Finding 3 – Sweet spot at 50 steps: FID is 45.11 at 50 steps, only 7 percent worse than optimal, at 20 times less computation than 1000 steps. This is the recommended deployment configuration.

Finding 4 – Steepest gains in early steps: Moving from 1 to 10 steps reduces FID by 86 percent. Moving from 10 to 50 steps reduces FID by a further 30 percent. Beyond 50 steps, gains become marginal.

C. Assignment 3 Experiment 2: DDIM Eta Parameter Study

TABLE VI
EXPERIMENT 2: FID vs. η (CIFAR-10, 10 STEPS FIXED, 1000 SAMPLES)

η	FID	Sampling Mode
0.00	65.72	Fully Deterministic (Best)
0.25	65.90	Mostly Deterministic
0.50	66.05	Mixed
0.75	72.54	Mostly Stochastic
1.00	89.55	Fully Stochastic (Worst)

Finding 1 – Deterministic is best: η equal to 0 achieves FID 65.72, confirming DDIM’s claim that deterministic trajectories outperform stochastic ones at constrained budgets.

Finding 2 – Stable zone up to η 0.5: FID changes by only 0.33 across η 0.0 to 0.5. Practitioners can safely use any η value up to 0.5 without meaningful quality loss.

Finding 3 – Phase transition above η 0.5: A 9.7 percent FID jump between η 0.5 and η 0.75, then a further 23 percent increase to η 1.0. This phase-transition behavior at η 0.5 was not described in the original DDIM paper.

Finding 4 – Full stochasticity costs 36 percent: FID 89.55 at η 1.0 versus 65.72 at η 0, with identical computation. This validates DDIM’s core contribution empirically on our setup.




Fig. 7. FID vs. eta at fixed 10 steps. Green zone (eta 0.0 to 0.5): stable region with under 0.5 percent FID change. Red zone (eta 0.5 to 1.0): sharp degradation where FID rises from 66.05 to 89.55. The phase transition at eta 0.5 is an original empirical finding of this work.

TABLE VII
EXPERIMENT 3: CROSS-DOMAIN CELEBA-HQ 256X256 RESULTS

Metric	Value
Model checkpoint	google/ddpm-celebahq-256
Resolution	256x256 pixels
Inference steps	1000 (full DDPM)
Batch size	4 (8GB VRAM constraint)
Time per image	Approximately 90 seconds
Total GPU time	Approximately 7 hours
Images generated	244
Intra-FID (diversity)	59.34
Literature FID range	29.76 to 40.26
Reference stats	Unavailable (HTTP 404)



Fig. 8. Generated CelebA-HQ 256x256 face images (4x4 grid from 244 generated). Diversity across identity, skin tone, gender, pose, and lighting confirms no mode collapse. Intra-FID of 59.34 quantifies this diversity.

D. Assignment 3 Experiment 3: Cross-Domain CelebA-HQ

Finding 1 – Cross-domain transfer confirmed: Generated images show realistic diverse human faces across varied identities and conditions. DDPM weights transfer meaningfully to an unseen face domain without fine-tuning.

Finding 2 – Diversity confirmed by intra-FID 59.34: Low intra-FID between two halves of the generated set confirms no mode collapse. Published FID for strong models on CelebA-HQ 256x256 ranges from 29.76 to 40.26, providing qualitative context.

Finding 3 – Infrastructure limitation: clean-fid returned HTTP 404 for CelebA-HQ reference stats. This is documented as a reproducibility infrastructure issue analogous to the Consistency Models checkpoint deletion.

E. Complete Results Summary

TABLE VIII
ALL RESULTS: ASSIGNMENT 2 BASELINES AND ASSIGNMENT 3 EXPERIMENTS

Configuration	Steps	Eta	FID	Source
DDPM (paper)	1000	1.0	3.17	[1]
DDPM (ours)	1000	1.0	20.91	A2
DDIM	10	0.0	38.69	A2
DDIM	1	0.0	462.11	A2
DDIM optimal	200	0.0	42.13	A3 Exp1
DDIM efficient	50	0.0	45.11	A3 Exp1
DDIM eta=0	10	0.00	65.72	A3 Exp2
DDIM eta=1	10	1.00	89.55	A3 Exp2
CelebA intra-FID	1000	1.0	59.34	A3 Exp3




Fig. 9. Comprehensive FID comparison across all A2 baselines and A3 experiments on CIFAR-10 (log scale). The optimal A3 DDIM-200 (FID 42.13) is competitive with A2 DDIM-10 (FID 38.69) but does not surpass full DDPM-1000 stochastic sampling (FID 20.91), confirming that inference optimization reduces but does not eliminate the gap.

VI. DISCUSSION

A. Why DDIM Degrades at 1000 Steps

The degradation from FID 42.13 at 200 steps to 60.55 at 1000 steps is counterintuitive since more denoising steps should allow finer resolution of the reverse process. The explanation lies in the numerical properties of the first-order Euler ODE integrator used by the DDIM scheduler. When eta equals 0, DDIM follows a deterministic probability flow ODE. Each

Euler step introduces a small local truncation error proportional to the step size. Over 1000 steps these errors accumulate into a systematic bias in the generated image distribution that FID successfully detects. Karras et al. [4] demonstrated that higher-order Heun solvers dramatically reduce this error. Our findings empirically confirm the discretization error hypothesis on the diffusers library implementation without any change to model weights.

B. The Eta Phase Transition at 0.5

The sharp FID jump between eta 0.5 (FID 66.05) and eta 0.75 (FID 72.54), a 9.7 percent increase, versus the near-flat region from eta 0.0 to 0.5 (0.5 percent change), has practical implications. Applications requiring output diversity can use eta up to 0.5 without meaningful quality cost. Values above 0.5 impose an increasingly severe penalty. This threshold was not previously characterized in the DDIM paper and represents an original finding with practical deployment value.

C. Analysis of the Reproduction FID Gap

The gap between reproduced FID (20.91) and the paper (3.17) is expected. Research has shown that FID computed with 5,000 samples systematically overestimates the true FID relative to 50,000 samples due to increased distributional estimate variance. Using 10,000 real reference images instead of 50,000 further affects the statistics. The qualitative grids show recognizable objects, the trend across configurations matches the paper, and all generation scripts ran without errors, confirming the reproduction is functionally correct.

D. Reflection on Real-World Challenges

This project provided direct experience with challenges common in applied deep learning but rarely discussed in coursework. The RTX 5050 incompatibility required researching PyTorch nightly builds. The deleted Consistency Models checkpoint reflects the growing problem of artifact deprecation where published peer-reviewed results cannot be verified because authors delete resources without notice. The power interruption demonstrated the importance of background job management and process state recovery using Linux signals. The duplicate process race condition taught strict process checking discipline. Each challenge was documented and resolved, contributing to a research record that reflects genuine engineering effort over several weeks.

E. Limitations

Sample size: All Assignment 3 CIFAR-10 FID estimates use 1,000 generated samples, carrying higher variance than the 5,000-sample Assignment 2 estimates. Generating 5,000 samples across 14 configurations was computationally infeasible on a laptop GPU.

Cross-domain evaluation: The 244-sample CelebA-HQ evaluation is insufficient for reliable absolute FID. Intra-FID provides directional diversity evidence but cannot be benchmarked against published results with different sample sizes.

No higher-order solvers: Verifying the discretization error hypothesis fully requires a second-order solver implementation, which is beyond this project’s scope.

No retraining: All experiments focus on inference-time hyperparameters. Training-side improvements such as EDM preconditioning [4] and cosine noise schedules [5] would likely further close the FID gap.

VII. CONCLUSION

This paper presented a complete account of a DDPM reproducibility study and inference-time optimization project spanning multiple weeks on a student laptop, involving generation of over 24,000 images across CIFAR-10 and CelebA-HQ, and resolution of seven distinct infrastructure and hardware challenges.

The core conclusions are as follows. The optimal DDIM step count for CIFAR-10 is 200 steps (FID 42.13), not the maximum 1000. FID degrades 44 percent beyond 200 steps due to first-order ODE discretization error accumulation, an empirical confirmation of the EDM hypothesis. For deployment prioritizing speed, 50 steps achieves 94 percent of optimal quality at 20 percent of computational cost. Deterministic sampling with eta equal to 0 outperforms full stochastic sampling by 36 percent at a fixed 10-step budget. A stable quality zone exists for eta from 0.0 to 0.5, with phase-transition-like degradation above 0.5 that was not previously characterized in the DDIM literature. Pre-trained DDPM weights transfer meaningfully to human face synthesis at 256x256 resolution, producing diverse outputs confirmed by intra-FID 59.34 after 7 hours of continuous laptop GPU computation.

Future work could apply higher-order Heun ODE solvers to eliminate the 1000-step degradation, conduct the eta study at multiple step budgets to map a full two-dimensional quality surface, extend cross-domain evaluation to LSUN and ImageNet with sufficient compute for absolute FID computation, and explore training-side improvements to close the reproduction FID gap.

REFERENCES

- [1] J. Ho, A. Jain, and P. Abbeel, “Denoising Diffusion Probabilistic Models,” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 6840–6851, 2020. arXiv:2006.11239. Code: <https://github.com/hojonathanho/diffusion>
- [2] J. Song, C. Meng, and S. Ermon, “Denoising Diffusion Implicit Models,” *International Conference on Learning Representations (ICLR)*, 2021. arXiv:2010.02502.
- [3] Y. Song, P. Dhariwal, M. Chen, and I. Sutskever, “Consistency Models,” *International Conference on Machine Learning (ICML)*, 2023. arXiv:2303.01469.
- [4] T. Karras, M. Aittala, T. Aila, and S. Laine, “Elucidating the Design Space of Diffusion-Based Generative Models,” *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2206.00364.
- [5] A. Q. Nichol and P. Dhariwal, “Improved Denoising Diffusion Probabilistic Models,” *International Conference on Machine Learning (ICML)*, 2021. arXiv:2102.09672.
- [6] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer, “High-Resolution Image Synthesis with Latent Diffusion Models,” *IEEE/CVF CVPR*, 2022. arXiv:2112.10752.
- [7] W. Peebles and S. Xie, “Scalable Diffusion Models with Transformers,” *IEEE/CVF ICCV*, 2023. arXiv:2212.09748.
- [8] D. Podell et al., “SDXL: Improving Latent Diffusion Models for High-Resolution Image Synthesis,” *ICLR*, 2024. arXiv:2307.01952.

- [9] J. Sohl-Dickstein, E. Weiss, N. Maheswaranathan, and S. Ganguli, "Deep Unsupervised Learning using Nonequilibrium Thermodynamics," ICML, 2015.
- [10] P. von Platen et al., "Diffusers: State-of-the-art Diffusion Models," HuggingFace, 2022. <https://github.com/huggingface/diffusers>
- [11] G. Parmar, R. Zhang, and J. Y. Zhu, "On Aliased Resizing and Surprising Subtleties in GAN Evaluation," IEEE/CVF CVPR, 2022. arXiv:2104.11222.
- [12] A. Vahdat, K. Kreis, and J. Kautz, "Score-based Generative Modeling in Latent Space," NeurIPS, 2021. arXiv:2106.05931.
- [13] F. A. Croitoru et al., "Diffusion Models in Vision: A Survey," IEEE TPAMI, 2023. arXiv:2209.04747.